

STOCHASTIC LOCAL SEARCH  
FOUNDATIONS AND APPLICATIONS

SLS Methods: An Overview

Holger H. Hoos & Thomas Stützle

# Outline

---

1. Iterative Improvement (Revisited)
2. 'Simple' SLS Methods
3. Hybrid SLS Methods
4. Population-based SLS Methods

In II, various mechanisms (*pivoting rules*) can be used for choosing improving neighbour in each step:

- ▶ *Best Improvement* (aka *gradient descent*, *greedy hill-climbing*): Choose maximally improving neighbour

*Note:* Requires evaluation of all neighbours in each step.

- ▶ *First Improvement*: Evaluate neighbours in fixed order, choose first improving step encountered.

*Note:* Can be much more efficient than Best Improvement; order of evaluation can have significant impact on performance.

# 'Simple' SLS Methods

---

Goal:

Effectively escape from local minima of given evaluation function.

General approach:

For fixed neighbourhood, use step function that permits *worsening search steps*.

Specific methods:

- ▶ Randomised Iterative Improvement
- ▶ Probabilistic Iterative Improvement
- ▶ Simulated Annealing
- ▶ Tabu Search
- ▶ Dynamic Local Search

## Randomised Iterative Improvement

**Key idea:** In each search step, with a fixed probability perform an uninformed random walk step instead of an iterative improvement step.

### Randomised Iterative Improvement (RII):

determine initial candidate solution  $s$

While termination condition is not satisfied:

    With probability  $wp$ :

        choose a neighbour  $s'$  of  $s$  uniformly at random

    Otherwise:

        choose a neighbour  $s'$  of  $s$  such that  $g(s') < g(s)$  or,

        if no such  $s'$  exists, choose  $s'$  such that  $g(s')$  is minimal

$s := s'$

## Note:

- ▶ No need to terminate search when local minimum is encountered  
*Instead:* Bound number of search steps or CPU time from beginning of search or after last improvement.
- ▶ Probabilistic mechanism permits arbitrary long sequences of random walk steps  
*Therefore:* When run sufficiently long, RII is guaranteed to find (optimal) solution to any problem instance with arbitrarily high probability.
- ▶ A variant of RII has successfully been applied to SAT (GWSAT algorithm), but generally, RII is often outperformed by more complex SLS methods.

## Probabilistic Iterative Improvement

**Key idea:** Accept worsening steps with probability that depends on respective deterioration in evaluation function value:  
bigger deterioration  $\cong$  smaller probability

*Realisation:*

- ▶ Function  $p(g, s)$ : determines probability distribution over neighbours of  $s$  based on their values under evaluation function  $g$ .
- ▶ Let  $step(s)(s') := p(g, s)(s')$ .

*Note:*

- ▶ Behaviour of PII crucially depends on choice of  $p$ .
- ▶ II and RII are special cases of PII.

## Simulated Annealing

**Key idea:** Vary temperature parameter, *i.e.*, probability of accepting worsening moves, in Probabilistic Iterative Improvement according to *annealing schedule* (aka *cooling schedule*).

### Inspired by physical annealing process:

- ▶ candidate solutions  $\cong$  states of physical system
- ▶ evaluation function  $\cong$  thermodynamic energy
- ▶ globally optimal solutions  $\cong$  ground states
- ▶ parameter  $T \cong$  physical temperature

*Note:* In physical process (e.g., annealing of metals), perfect ground states are achieved by very slow lowering of temperature.



## Simulated Annealing (SA):

determine initial candidate solution  $s$

set initial temperature  $T$  according to *annealing schedule*

While termination condition is not satisfied:

    | probabilistically choose a neighbour  $s'$  of  $s$   
    |     using *proposal mechanism*

    | If  $s'$  satisfies probabilistic *acceptance criterion* (depending on  $T$ ):

$s := s'$

    | update  $T$  according to *annealing schedule*

## Note:

- ▶ 2-stage step function based on
  - ▶ proposal mechanism (often uniform random choice from  $N(s)$ )
  - ▶ acceptance criterion (often *Metropolis condition*)
- ▶ Annealing schedule (function mapping run-time  $t$  onto temperature  $T(t)$ ):
  - ▶ initial temperature  $T_0$   
(may depend on properties of given problem instance)
  - ▶ temperature update scheme  
(e.g., geometric cooling:  $T := \alpha \cdot T$ )
  - ▶ number of search steps to be performed at each temperature  
(often multiple of neighbourhood size)
- ▶ Termination predicate: often based on *acceptance ratio*, i.e., ratio of proposed vs accepted steps.

## 'Convergence' result for SA:

Under certain conditions (extremely slow cooling), any sufficiently long trajectory of SA is guaranteed to end in an optimal solution [Geman and Geman, 1984; Hajek, 1998].

### Note:

- ▶ Practical relevance for combinatorial problem solving is very limited (impractical nature of necessary conditions)
- ▶ In combinatorial problem solving, *ending* in optimal solution is typically unimportant, but *finding* optimal solution during the search is (even if it is encountered only once)!

## Tabu Search

**Key idea:** Use aspects of search history (memory) to escape from local minima.

### Simple Tabu Search:

- ▶ Associate *tabu attributes* with candidate solutions or solution components.
- ▶ Forbid steps to search positions recently visited by underlying iterative best improvement procedure based on tabu attributes.

## Tabu Search (TS):

determine initial candidate solution  $s$

While *termination criterion* is not satisfied:

| *determine set  $N'$  of non-tabu neighbours of  $s$*   
| *choose a best improving candidate solution  $s'$  in  $N'$*   
| *update tabu attributes* based on  $s'$   
|  $s := s'$

## Note:

- ▶ Non-tabu search positions in  $N(s)$  are called *admissible neighbours of  $s$* .
- ▶ After a search step, the current search position or the solution components just added/removed from it are declared *tabu* for a fixed number of subsequent search steps (*tabu tenure*).
- ▶ Often, an additional *aspiration criterion* is used: this specifies conditions under which tabu status may be overridden (e.g., if considered step leads to improvement in incumbent solution).

**Note:** Performance of Tabu Search depends crucially on setting of tabu tenure  $tt$ :

- ▶  $tt$  too low  $\Rightarrow$  search stagnates due to inability to escape from local minima;
- ▶  $tt$  too high  $\Rightarrow$  search becomes ineffective due to overly restricted search path (admissible neighbourhoods too small)

Further improvements can be achieved by using *intermediate-term* or *long-term memory* to achieve additional *intensification* or *diversification*.

### Examples:

- ▶ Occasionally backtrack to *elite candidate solutions*, i.e., high-quality search positions encountered earlier in the search; when doing this, all associated tabu attributes are cleared.
- ▶ Freeze certain solution components and keep them fixed for long periods of the search.
- ▶ Occasionally force rarely used solution components to be introduced into current candidate solution.
- ▶ Extend evaluation function to capture frequency of use of candidate solutions or solution components.



Tabu search algorithms are state of the art for solving many combinatorial problems, including:

- ▶ SAT and MAX-SAT
- ▶ the Constraint Satisfaction Problem (CSP)
- ▶ many scheduling problems

Crucial factors in many applications:

- ▶ choice of neighbourhood relation
- ▶ efficient evaluation of candidate solutions  
(caching and incremental updating mechanisms)

# Hybrid SLS Methods

---

Combination of 'simple' SLS methods often yields substantial performance improvements.

Simple examples:

- ▶ Commonly used restart mechanisms can be seen as hybridisations with Uninformed Random Picking
- ▶ Iterative Improvement + Uninformed Random Walk = Randomised Iterative Improvement

## Iterated Local Search

**Key Idea:** Use two types of SLS steps:

- ▶ *subsidiary local search* steps for reaching local optima as efficiently as possible (intensification)
- ▶ *perturbation steps* for effectively escaping from local optima (diversification).

Also: Use *acceptance criterion* to control diversification vs intensification behaviour.

## Iterated Local Search (ILS):

determine initial candidate solution  $s$

perform *subsidiary local search* on  $s$

While termination criterion is not satisfied:

$r := s$

    perform *perturbation* on  $s$

    perform *subsidiary local search* on  $s$

    based on *acceptance criterion*,

        keep  $s$  or revert to  $s := r$

## Note:

- ▶ *Subsidiary local search* results in a local minimum.
- ▶ ILS trajectories can be seen as walks in the space of local minima of the given evaluation function.
- ▶ *Perturbation phase* and *acceptance criterion* may use aspects of *search history* (i.e., limited memory).
- ▶ In a high-performance ILS algorithm, *subsidiary local search*, *perturbation mechanism* and *acceptance criterion* need to complement each other well.

## Subsidiary local search:

- ▶ More effective subsidiary local search procedures lead to better ILS performance.

*Example: 2-opt vs 3-opt vs LK for TSP.*

- ▶ Often, subsidiary local search = iterative improvement, but more sophisticated SLS methods can be used. (e.g., Tabu Search).

## Perturbation mechanism:

- ▶ Needs to be chosen such that its effect *cannot* be easily undone by subsequent local search phase.  
(Often achieved by search steps larger neighbourhood.)

*Example:* local search = 3-opt, perturbation = 4-exchange steps in ILS for TSP.

- ▶ A perturbation phase may consist of one or more perturbation steps.

## Perturbation mechanism (continued):

- ▶ Weak perturbation  $\Rightarrow$  short subsequent local search phase;  
*but*: risk of revisiting current local minimum.
- ▶ Strong perturbation  $\Rightarrow$  more effective escape from local minima; *but*: may have similar drawbacks as random restart.
- ▶ Advanced ILS algorithms may change nature and/or strength of perturbation adaptively during search.



## Acceptance criteria:

- ▶ Always accept the *better* of the two candidate solutions  
⇒ ILS performs Iterative Improvement in the space of local optima reached by subsidiary local search.
- ▶ Always accept the *more recent* of the two candidate solutions  
⇒ ILS performs random walk in the space of local optima reached by subsidiary local search.
- ▶ Intermediate behaviour: select between the two candidate solutions based on the *Metropolis criterion* (e.g., used in *Large Step Markov Chains* [Martin *et al.*, 1991]).
- ▶ Advanced acceptance criteria take into account search history, e.g., by occasionally reverting to *incumbent solution*.

## Iterated local search algorithms . . .

- ▶ are typically rather easy to implement (given existing implementation of subsidiary simple SLS algorithms);
- ▶ achieve state-of-the-art performance on many combinatorial problems, including the TSP.

# Population-based SLS Methods

---

SLS methods discussed so far manipulate one candidate solution of given problem instance in each search step.

**Straightforward extension:** Use *population* (i.e., set) of candidate solutions instead.

## Note:

- ▶ The use of populations provides a generic way to achieve search diversification.
- ▶ Population-based SLS methods fit into the general definition from Chapter 1 by treating sets of candidate solutions as search positions.